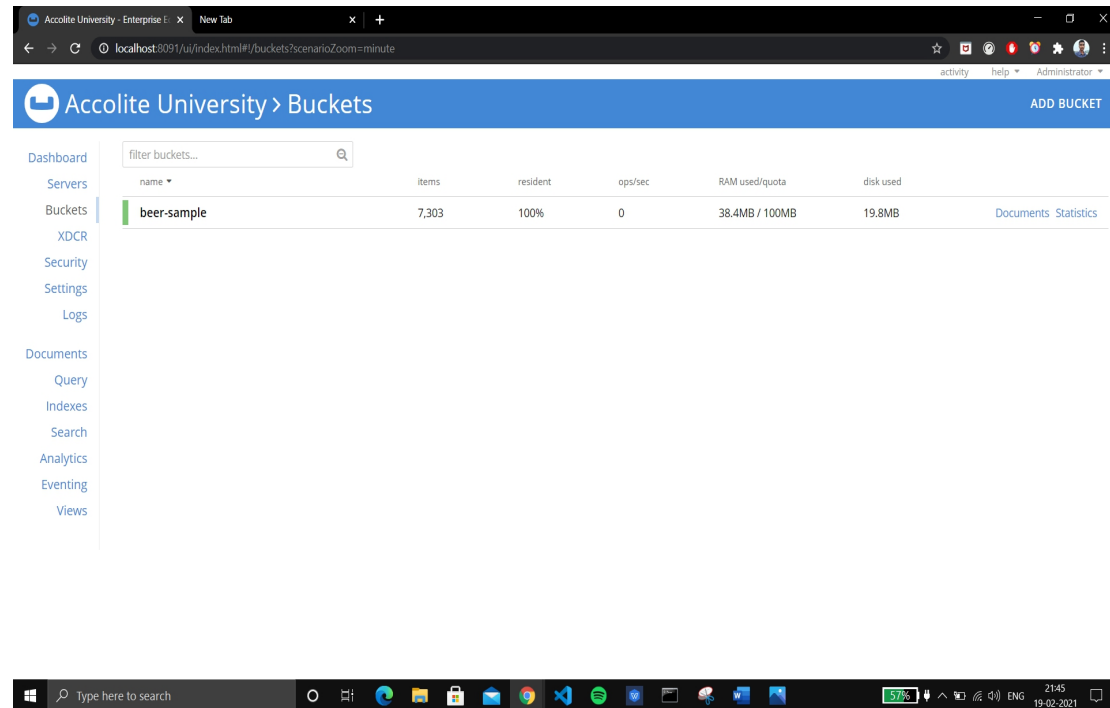


NO SQL Assignment

1. Import `beer-sample` bucket



2. Write a join query to fetch the Top 10 brewery(type="beer") and their country(type="brewery") which produces more varieties of beers.

```
select b1.name as beer_name,  
b2.country as country_name  
FROM `beer-sample` b1 JOIN `beer-sample` b2 ON KEYS b1.brewery_id  
WHERE b1.type="beer"  
AND b2.type="brewery"  
LIMIT 10
```

Output:

The screenshot shows the Accolite University Query Editor interface. The query editor contains the following SQL query:

```
1 select b1.name as beer_name,
2 b2.country as country_name
3 FROM `beer-sample` b1 JOIN `beer-sample` b2 ON KEYS b1.brewery_id
4 WHERE b1.type="beer"
5 AND b2.type="brewery"
6 LIMIT 10
```

The query results are displayed in a table format:

beer_name	country_name
21A IPA	United States
563 Stout	United States
Amendment Pale Ale	United States
Bitter American	United States
Double Trouble IPA	United States
General Pippo's Porter	United States
North Star Red	United States
Oyster Point Oyster Stout	United States
Potrero ESB	United States
South Park Blonde	United States

The interface also includes a sidebar with navigation options (Dashboard, Servers, Buckets, XDCR, Security, Settings, Logs, Documents, Query, Indexes, Search, Analytics, Eventing, Views) and a right-hand panel for Data Insights.

beer_name country_name
 "21A IPA" "United States"
 "563 Stout" "United States"
 "Amendment Pale Ale" "United States"
 "Bitter American" "United States"
 "Double Trouble IPA" "United States"
 "General Pippo's Porter" "United States"
 "North Star Red" "United States"
 "Oyster Point Oyster Stout" "United States"
 "Potrero ESB" "United States"
 "South Park Blonde" "United States"

3. Write a mapreduce to get the number of breweries based on country.
 Please attach the mapreduce code and json output screenshot.

The screenshot shows the Accolite University MapReduce interface. The Map code is as follows:

```
1 function (doc, meta) {
2   if(doc.type=="brewery")
3     emit(doc.country, null);
4 }
```

The Reduce code is as follows:

```
1 _count
```

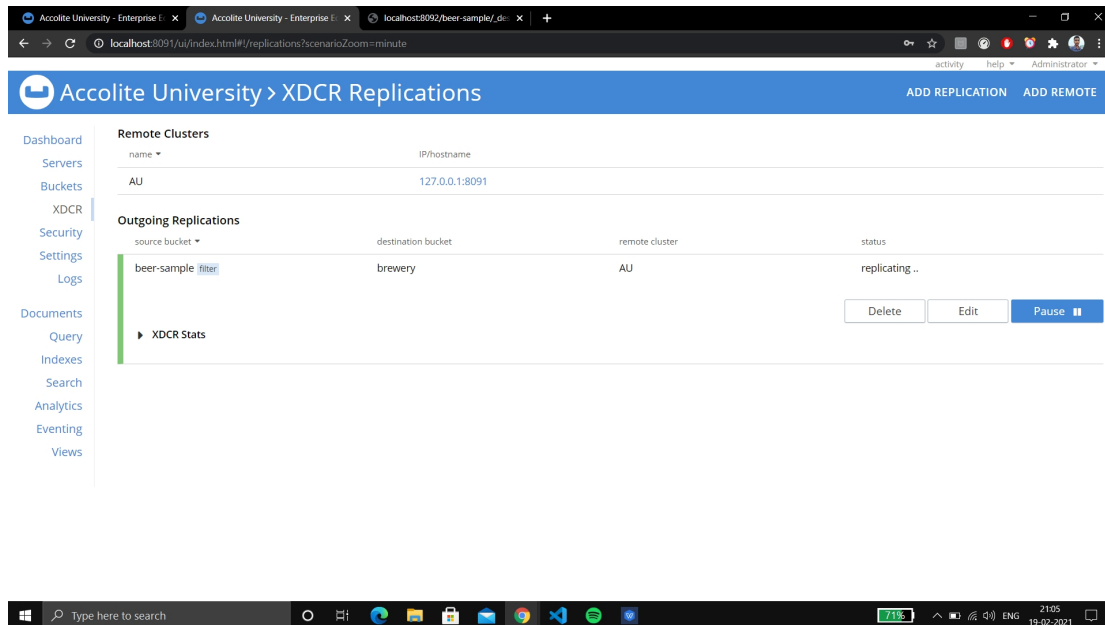
The results are displayed in a table format:

Key	Value
null	1412
undefined	

The interface also includes a sidebar with navigation options (Documents, Query, Indexes, Search, Analytics, Eventing, Views) and a right-hand panel for Data Insights.

4. XDCR:

- Add a new bucket "Brewery".
- Create a XDCR with a filter(type='brewery') to replicate only the brewery entity from 'beer-sample' bucket.



5. CLI:

- Add a new bucket "Beer".
- Using CLI - do a cbexport of the entire 'beer-sample'
- And do a cbimport with "brewery_id" as primary key. As a result, in the new bucket - only "beer" documents will be imported with their respective brewery name as meta().id
- Share your observation on counts with screenshot.

